

Synchronizing Physical Clock Logical Clock & Vector Clock

Y. V. Lokeswari

AP/CSE

Reference: 1. Mukesh Singhal & N.G. Shivaratri, Advanced
Concepts in Operating Systems,

2. George Coulouris, Jean Dollimore and Tim Kindberg,
“Distributed Systems Concepts and Design”, Fifth Edition,
Pearson Education, 2012



Overview

- Physical Clocks
- Synchronizing Physical Clock (Algorithms)
- Problems with Physical Clock
- Lamport's Logical Clock
- Problems with Logical Clock
- Vector Clock




Clock Synchronization

- Temporal ordering of events produced by concurrent processes
- Synchronization between senders and receivers of messages
- Serialization of concurrent access for shared objects
- **Physical Clock:** It is the internal clock present in a computer.
(Time of a day)
- **Logical Clock:** keeps track of event ordering among related
(causal) events.



Clock Synchronization

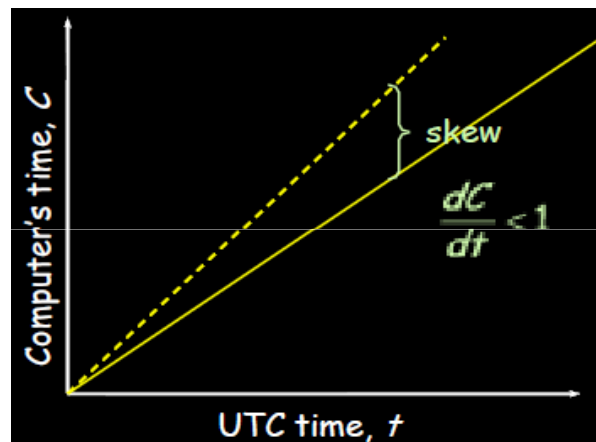

- Getting two systems to agree on time
 - Two clocks hardly ever agree
 - **Quartz oscillators oscillate at slightly different frequencies**
- **Clock Drift** : Clocks tick at different rates. Create ever-widening gap in perceived time
 - Clock drift in ordinary clocks based on quartz crystal is 10^{-6} seconds.
 - This creates a difference of 1 sec for every 11.6 days (1,000,000 sec)
 - Clock drift of high precision clock is 10^{-7} to 10^{-8}
- **Clock Skew** : Difference between two clocks at one point in time.



Clock Synchronization


- Dealing with Clock Drift: Go for gradual clock correction

UTC : Universal Coordinated Time



If slow:

Make clock run faster until
it Synchronizes



If fast:

Make clock run slower
until it Synchronizes



Clock Synchronization

- **External Synchronization** : Clock synchronizes to correct time from external timing elements like Radio / Satellite time.
- **Internal Synchronization** : Clock synchronizes to correct time by getting timings from other computers.
- 3 - Algorithms for Synchronizing Physical Clock
 1. Cristian's Algorithm
 2. Berkeley Algorithm
 3. Network Time Protocol (NTP)




Clock Synchronization

1. Cristian's Algorithm

$(T_1 - T_0)/2$ is round-trip time

Client sets time to:



$$T_{new} = T_{server} + \frac{T_1 - T_0}{2}$$

Clock Synchronization

1. Cristian's Algorithm – Example

- Send request at 5:08:15.100 (T₀)
- Receive response at 5:08:15.900 (T₁)
- Response contains 5:09:25.300 (T_{server})
- Elapsed time is T₁ - T₀

$$5:08:15.900 - 5:08:15.100 = 800 \text{ msec}$$

- Best guess: timestamp was generated
400 msec ago (800/2)
- Set time to T_{server} + elapsed time

$$5:09:25.300 + 400 = 5:09:25.700$$



Clock Synchronization

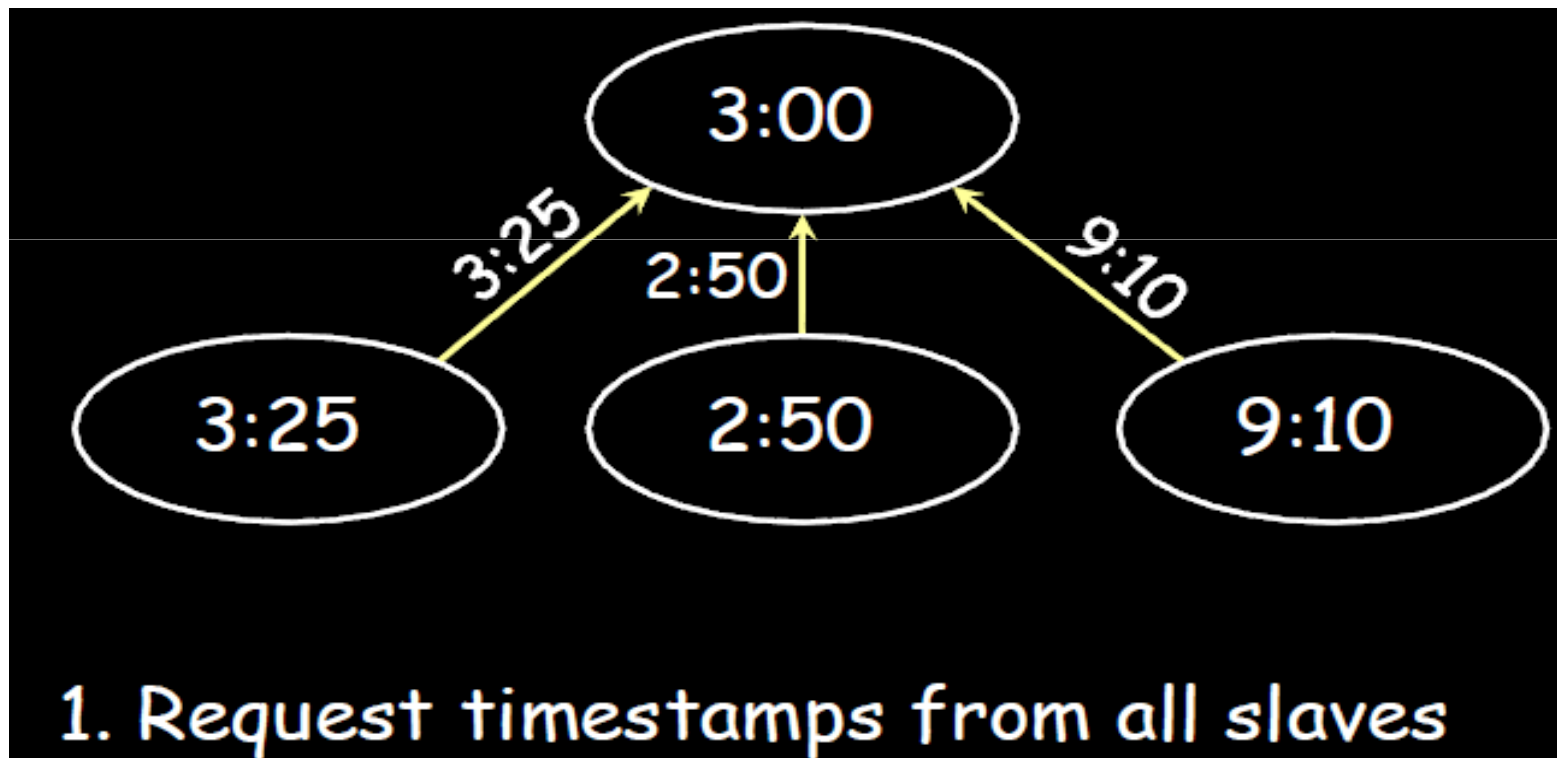
2. Berkeley Algorithm

- Machines run **time daemon** Process that implements protocol
- One machine is elected (or designated) as the server (**master**) Others are **slaves**
- Master polls each machine periodically and ask each machine for time
- Can use Cristian's algorithm to compensate for network latency
- When results are received, master computes average of times - Including master's time
- **Hope**: Average cancels out individual clock's tendencies to run fast or slow (clock drift)
- Send offset by which each clock needs adjustment to each slave.
- Algorithm has provisions for ignoring readings from clocks whose skew is too great
- Compute a **fault-tolerant average**
- If master fails - any slave can take over




Clock Synchronization

2. Berkeley Algorithm Example



Clock Synchronization

2. Berkeley Algorithm Example




2. Compute fault-tolerant average:

$$\frac{3:25 + 2:50 + 3:00}{3} = 3:05$$

Clock Synchronization

2. Berkeley Algorithm Example



3. Send offset to each client

Clock Synchronization

3. Network Time Protocol


- Enable clients across Internet to be accurately synchronized to UTC despite message delays
- Primary servers are connected directly to a time source such as a radio clock receiving UTC; secondary servers are synchronized with primary servers.
- The servers are connected in a logical hierarchy called a **synchronization subnet** whose levels are called **strata**.
- Primary servers occupy stratum 1: they are at the root.
- Stratum 2 servers are secondary servers that are synchronized directly with the primary servers;
- Stratum 3 servers are synchronized with stratum 2 servers, and so on.



Clock Synchronization

3. Network Time Protocol

Arrows denote synchronization control, numbers denote strata.



Clock Synchronization

3. Network Time Protocol (Synchronization Modes)

- **Multicast mode**
 - Every node sends its time to all the other nodes in the LAN
 - For high speed LANs
 - Lower accuracy but efficient
- **Procedure call mode**
 - Similar to Cristian's algorithm
- **Symmetric mode**
 - Intended for master servers
 - Pair of servers exchange messages and retain data to improve synchronization over time

All messages delivered unreliably with UDP



Problems with Physical Clocks

- **Clock Drift** : Clocks tick at different rates. Create ever-widening gap in perceived time
- **Clock Skew** : Difference between two clocks at one point in time.
- For quartz crystal clocks, typical drift rate is about one second every 10^6 seconds ≈ 11.6 days
- Best atomic clocks have drift rate of one second in 10^{13} seconds $\approx 300,000$ years.
- **Quartz clock run at rate of 1.5 microseconds slower for every 35 days.**



Problems with Physical Clocks



8:01:24

Skew = +84 seconds
+84 seconds/35 days
Drift = +2.4 sec/day

Oct 23, 2006
8:00:00



8:01:48

Skew = +108 seconds
+108 seconds/35 days
Drift = +3.1 sec/day



Logical Clocks



Need for Logical Clock


- For many purposes, it is sufficient to know the order in which events occurred.
- Lamport (1978) — introduce logical (*virtual*) time, synchronize *logical clocks*.
- An event may be an instruction execution, may be a function execution, etc.
- Events include message send / receive

Within a single process, or between two processes on the same computer

- The order in which two **events** occur **can** be determined using the **physical clock**

Between two different computers in a distributed system

- The order in which two events occur **cannot** be determined **using local physical clocks**, since those clocks cannot be synchronized perfectly



Happened Before Relation

- Lamport defined the happened before relation (denoted as “ \rightarrow ”), which describes a **causal ordering of events**:
 1. if a and b are events in the same process, and a occurred before b, then $a \rightarrow b$
 2. if a is the event of sending a message m in one process, and b is the event of receiving that message m in another process, then $a \rightarrow b$
 3. if $a \rightarrow b$, and $b \rightarrow c$, then $a \rightarrow c$ (i.e., the relation “ \rightarrow ” is transitive)




Causality

- Past events influence future events
- This influence among causally related events (those that can be ordered by “ \rightarrow ”) is referred to a **causally affects**.
- If $a \rightarrow b$, event **a** causally affects event **b**


Concurrent events

- Two distinct events **a** and **b** are said to be concurrent (denoted “ $a \parallel b$ ”), if neither $a \rightarrow b$ nor $b \rightarrow a$
- In other words, concurrent events do not causally affect each other
- For any two events a and b in a system, either: $a \rightarrow b$ or $b \rightarrow a$ or $a \parallel b$



Causal Events


Space-time diagram



Concurrent Events

Space-time diagram

events in the set $\{e_1^3, e_2^4, e_3^3\}$ are logically concurrent.



Lamport's Logical Clock

- To implement “ \rightarrow ” in a distributed system, Lamport (1978) introduced the concept of logical clocks, which captures “ \rightarrow ” numerically
- Each process P_i has a logical clock C_i
- Clock C_i can assign a value $C_i(a)$ to any event **a** in process P_i
- The value $C_i(a)$ is called the timestamp of event **a** in process P_i
- The value $C(a)$ is called the timestamp of event **a** in whatever process it occurred.
- The timestamps have no relation to physical time, which leads to the term logical clock.
- The logical clocks assign monotonically increasing timestamps, and can be implemented by simple counters



Lamport's Logical Clock

- **Clock condition:** if $a \rightarrow b$, then $C(a) < C(b)$
- If event a happens before event b , then the clock value (timestamp) of a should be less than the clock value of b
- Note that we **can not say: if $C(a) < C(b)$, then $a \rightarrow b$**
- **Correctness conditions (must be satisfied by the logical clocks to meet the clock condition above):**
- [C1] For any two events a and b in the same process P_i ,
if a happens before b , then $C_i(a) < C_i(b)$
- [C2] If event a is the event of sending a message m in process P_i , and event b is the event of receiving that same message m in a different process P_k , then
 $C_i(a) < C_k(b)$



Lamport's Logical Clock

- **[IR1]** Clock C_i must be incremented between any two successive events in process P_i

$$C_i := C_i + d, \quad (d > 0) \text{ (usually } d=1\text{)}$$

- **[IR2]** If event a is the event of sending a message m in process P_i , then message m is assigned a timestamp $tm = C_i(a)$. When that same message m is received by a different process P_k , C_k is set to a value greater than or equal to its present value, and greater than tm .


$$C_k := \max(C_k, tm + d), \quad (d > 0) \text{ (usually } d=1\text{)}$$



Lamport's Logical Clock

IR 1: $C_i := C_i + d$, ($d > 0$), (usually $d=1$)


IR2: $C_k := \max(C_k, t_m + d)$, ($d > 0$) (Usually $d=1$)



Lamport's Logical Clock

IR 1: $C_i := C_i + d$, ($d > 0$), (usually $d=1$)


IR2: $C_k := \max(C_k, t_m + d)$, ($d > 0$) (Usually $d=1$)




Lamport's Logical Clock

- A total order of events (“ \Rightarrow ”) can be obtained as follows:
- If a is any event in process P_i , and b is any event in process P_k , then $a \Rightarrow b$ if and only if either:
- $C_i(a) < C_k(b)$ or
- $C_i(a) = C_k(b)$ and $P_i << P_k$

where “ $<<$ ” denotes a relation that totally orders the processes to break ties



Limitations of Lamport's Logical Clock



- With Lamport's logical clocks,
- if $a \rightarrow b$, then $C(a) < C(b)$
- The following is **not necessarily true if events a and b occur in** different processes:

- if $C(a) < C(b)$, then $a \rightarrow b$

$C(e12) < C(e23)$, and $e12 \rightarrow e23$ is true

$C(e12) < C(e24)$, but $e12 \rightarrow e24$ is false

$C(e15) < C(e25)$, but $e15 \rightarrow e25$ is false

- Cannot determine whether two events are causally related from timestamps



Vector Clocks



Vector Clock

- Maintain a vector of values for every event that happens in all processes.
- Update happens for group of values in every event.



Vector Clock

- [IR1] Clock C_i must be incremented between any two successive events in process P_i :

$$C_i[i] := C_i[i] + d, \quad (d > 0, \text{ usually } d=1)$$

- [IR2] If event a is the event of sending a message m in process P_i , then message m is assigned a vector timestamp $tm = C_i(a)$ When that same message m is received by a different process P_k , C_k is updated as follows:

$$\forall p, C_k[p] := \max(C_k[p], tm[p] + d), \quad (\text{usually } d=0 \text{ unless needed to model network delay})$$


- It can be shown that $\forall i, \forall k : C_i[i] \geq C_k[i]$
- Rules for comparing timestamps can also be established so that
- if $ta < tb$, then $a \rightarrow b$ / Solves the problem with Lamport's clocks



Vector Clock

IR1: $C_i[i] := C_i[i] + d$. ($d=1$)

IR2: $\forall p, C_k[p] := \max(C_k[p], tm[p] + d)$. ($d=0$)




“ enn ” is
event;
“ (n,n,n) ” is
clock value

Vector Clock – Solving Logical Clock Problem

if $t(e_{31}) < t(e_{23})$ then $e_{31} \rightarrow e_{23}$ is true.

if $t(e_{12}) < t(e_{24})$ then $e_{12} \rightarrow e_{24}$ is also true.

Concurrent events $e_{32} \parallel e_{24}$; $e_{32} \not\rightarrow e_{24}$ and $e_{24} \not\rightarrow e_{32}$



Vector Clock

Example of Vector Clock




Figure 4.3: Evolution of vector time.

Summary

- Physical Clocks
- Synchronizing Physical Clock (Algorithms)
- Problems with Physical Clock
- Lamport's Logical Clock
- Problems with Logical Clock
- Vector Clock





Thank You

